

Threads

ECE2893

Lecture 13

Threads

- 1 In order to implement concurrent algorithms, such as the parallel bubble sort discussed previously, we need some way to say that we want more than one CPU executing our program.
- 2 This is done by creating one or more *threads*.
- 3 A *thread* is nothing more than a sequence of C/C++ instructions.
- 4 All of our programs to date in fact use a thread!
 - ▶ All C/C++ programs start out with exactly one thread.
 - ▶ The `main` function is the starting point for the thread.
- 5 More threads can be created, using the *pthread*s library.
- 6 When creating a thread, the *starting function* is specified, and any arguments for that function.
- 7 The details of the *pthread*s library are quite complex and tedious.
- 8 We will be using a simplified, and reduced functionality, version called *gthreads*.

DISCLAIMER!

- 1 The *gthreads* library is not part of the standard C or C++ library.
- 2 It was created by the ECE2893 instructor to simplify the way that threads are created and managed.
- 3 The definition of the *gthreads* functions are in `gthreads.h`
- 4 The implementation of the *gthreads* functions are in `gthreads.cc`
- 5 The *gthreads* library has significantly reduced functionality versus the *pthread*s library.
- 6 However, *gthreads* has sufficient functionality for the ECE2893 class assignments.

Creating a Thread using *gthreads*

- 1 A program can ask for another thread of execution by using the `CreateThread` function defined by the *gthreads* library.
- 2 The first argument to `CreateThread` is required, and is the name of the starting function for the new thread.
 - ▶ The starting function can be any function already defined in the C/C++ program.
- 3 Following the function name argument, there can be up to four more arguments to `CreateThread`, of any type.
 - ▶ These arguments are passed by value to the thread starting function when the thread starts executing.
 - ▶ The types of the arguments *must* match the types of the arguments in the thread starting function.
 - ▶ If they don't match, a compiler error occurs.
- 4 See the next slide for an example.

CreateThread Example

```
// Illustrate using CreateThread in gthreads
// George F. Riley, Georgia Tech, ECE2893, Spring 2011

#include "gthread.h"    // Must be included to use the gthreads library

void BubbleSort(int* d, int startingPoint, int length)
{ // This is the thread starting point
  // This is where, in this example, the sorting of array d will be done
}

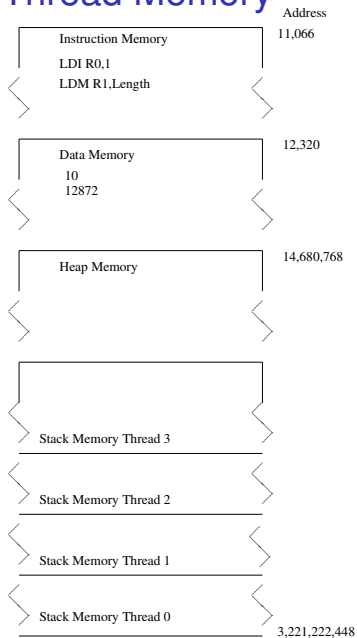
const int nThreads = 4;    // Number of threads desired
const int maxSize = 512000; // Largest sort size

int main()
{
  int d[maxSize];    // Array to be sorted
  int start = 0;    // Starting point of sub-array
  int length = (omitted); // Length of sub-array
  for (int k = 0; k < nThreads; ++k)
  { // Create each of the four sorting threads
    CreateThread(BubbleSort, d, start, length);
    // need more code here..(omitted)
  }
  // At this point all threads are created
  // Need more code here (shown later)
}
```

Thread Memory Use

- 1 The previous example created 4 threads, each starting at function `BubbleSort`, and passing three parameters to `BubbleSort`.
- 2 The 4 threads (in this example) have access to exactly the same memory (with one exception).
 - ▶ The thread can call any function that exists in the program.
 - ▶ The thread can read or write any global variable that exists in the program.
 - ▶ However, *EACH THREAD HAS A PRIVATE STACK*
 - ▶ This means that local variables within subroutines are not shared between threads.

Thread Memory



Thread Scheduling

- ❶ Calling `CreateThread` does *NOT* necessarily cause the `BubbleSort` function to be immediately executed!
- ❷ The operating system is free to assign CPU's to threads in any order, and switch CPU's between the threads.
- ❸ There can be (and often are) more threads than there are CPUs.
 - ▶ In this case, the operating system switches the CPU between the threads, just like switching between processes.
- ❹ The obvious question now is, “How do we delete a thread?”.
 - ▶ All threads stay alive until the thread starting function (`BubbleSort` in this example) exits.
 - ▶ When the thread function exits, the operating system deletes the thread and re-assigns the CPU to another thread or process.
 - ▶ **IMPORTANT.** If a threads “parent” (the thread creating the thread) exits, then all child threads are immediately terminated.
 - ▶ This means that when `main` exits (in the previous example), all of our `BubbleSort` threads terminate immediately, regardless of whether they are completed (or even started for that matter) or not.
 - ▶ We need some way for the `main` thread to wait for the completion of the child threads.

Thread Synchronization

- ➊ In *gthreads*, there are two co-operating functions that are used to let a parent function(the `main` function in the example) wait until all threads have finished processing.
- ➋ All threads *MUST* call `EndThread()` immediately before exiting.
 - ▶ This does several things, but primarily decrements a counter indicating how many active threads are remaining.
 - ▶ If this counter decrements to zero, the thread also *notifies* the parent function that all threads have finished.
- ➌ After all threads are created, the parent function *MUST* call `WaitAllThreads()` ;
 - ▶ This call tells the operating system that the `main` thread has no useful work to do until all of the child threads have completed and called `EndThread()`.
 - ▶ No CPU is assigned to `main` until all threads have finished.
- ➍ When `WaitAllThreads()` returns, the parent function is free to continue processing. In our case, it should perform the merging of the four sub-arrays into the final sorted array.

CreateThread Example, Better

```
// Illustrate using CreateThread in gthreads
// George F. Riley, Georgia Tech, ECE2893, Spring 2011

#include "gthread.h"    // Must be included to use the gthreads library

void BubbleSort(int* d, int startingPoint, int length)
{ // This is the thread starting point
  // This is where, in this example, the sorting of array d will be done
  EndThread(); // Call this just before exiting
}

const int nThreads = 4;    // Number of threads desired
const int maxSize = 512000; // Largest sort size

int main()
{
  int d[maxSize];          // Array to be sorted
  int start = 0;           // Starting point of sub-array
  int length = (omitted);  // Length of sub-array
  for (int k = 0; k < nThreads; ++k)
  { // Create each of the four sorting threads
    CreateThread(BubbleSort, d, start, length);
    // need more code here..(omitted)
  }
  // At this point all threads are created
  WaitAllThreads(); // This waits until all child threads are done
  // Perform the merge procedure to merge the separate sub-arrays
}
```